

# Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network

Group: C6 | Summer 2019

# Group Members

Md.Faisal Mahmud Abid (Id : 011 143 108)

Nopa Islam (Id : 011 151 062)

Nasir Uddin Ahmed (Id : 011 151 086)

# Paper Introduction

- Author :
  - a. Christian Ledig*
  - b. Lucas Theis*
  - c. Ferenc Huszar*
- *Number Of Citation : 1823*
- *Paper Presented On : Computer Vision and Pattern Recognition Conference (CVPR) 2017*

# Dataset

- Name : MIRFLICKR-25000.
- Consists of 25000 images.
- Pre-trained neural network


# Abstract

- Main problem; recovering finer texture details at large upscale factors.
- The paper proposes a generative adversarial network (GAN).
- Proposing a perceptual loss function.

# Introduction

- High **upscale** factor is a bigger issue.
- Texture details typically absent in reconstructed image.
- Minimizing the mean squared error (MSE).
- Minimizing (MSE) **maximizes** peak signal to noise ratio (PSNR).

# Basic Neural Network





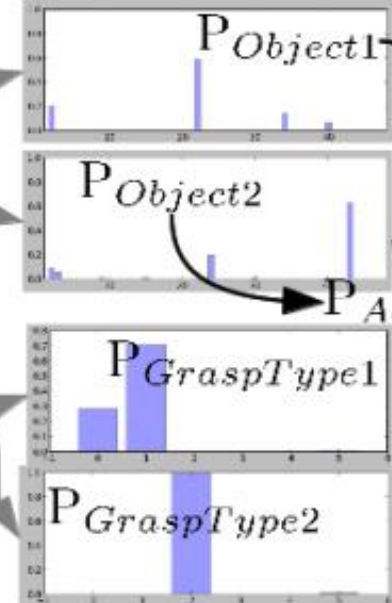

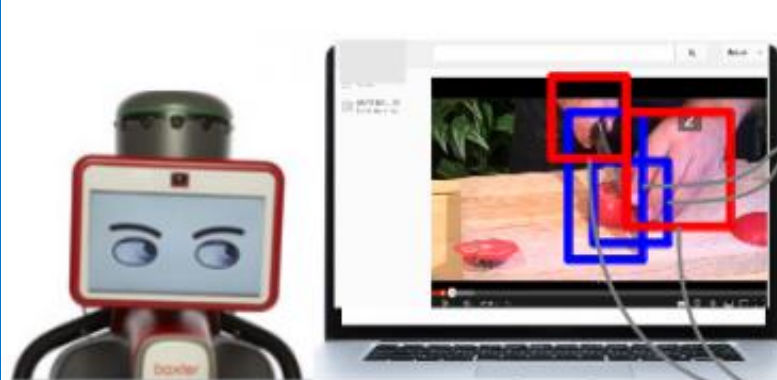


Figure : Application Of CNN





Most Likely Parse Tree

Viterbi


$AP$	$\rightarrow$	$G_1 O_1 \mid G_2 O_2 \mid A O_2 \mid A HP$	0.25	(1)
$HP$	$\rightarrow$	$H AP \mid HP AP$	0.5	(2)
$H$	$\rightarrow$	$'LeftHand' \mid 'RightHand'$	0.5	(3)
$G_1$	$\rightarrow$	$'GraspType1'$	$P_{GraspType1}$	(4)
$G_2$	$\rightarrow$	$'GraspType2'$	$P_{GraspType2}$	(5)
$O_1$	$\rightarrow$	$'Object1'$	$P_{Object1}$	(6)
$O_2$	$\rightarrow$	$'Object2'$	$P_{Object2}$	(7)
$A$	$\rightarrow$	$'Action'$	$P_{Action}$	(8)

# Design Of Convolutional Neural Network

# A toy ConvNet: X's and O's


Says whether a picture is of an X or an O

A two-dimensional  
array of pixels




**X** or **O**


# For example




# Trickier cases




translation



scaling







rotation



weight




X




O

# Deciding is hard



?

=












# ConvNets match pieces of the image



# Features match pieces of the image

1	-1	-1
-1	1	-1
-1	-1	1


1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1

1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1


-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1


-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1


-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1



-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1





# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

[illegible]

# Filtering: The math behind the match

1. Line up the feature and the image patch.
2. Multiply each image pixel by the corresponding feature pixel.
3. Add them up.
4. Divide by the total number of pixels in the feature.



# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{array}{|c|} \hline -1 \\ \hline \end{array} \times \begin{array}{|c|} \hline -1 \\ \hline \end{array} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{array}{|c|} \hline -1 \\ \hline \end{array} \times \begin{array}{|c|} \hline -1 \\ \hline \end{array} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{array}{|c|} \hline -1 \\ \hline \end{array} \times \begin{array}{|c|} \hline -1 \\ \hline \end{array} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1



Diagram illustrating the multiplication of two 1x1 matrices, resulting in a 1x1 matrix. The first 1x1 matrix is highlighted with a red box, and the second 1x1 matrix is highlighted with a yellow box. The result is a 1x1 matrix, also highlighted with a yellow box.

A 9x9 grid representing a sparse matrix  $A$ . The diagonal elements are all 1, while all other elements are -1. A teal square highlights the top-left 4x4 submatrix. Within this, a yellow square highlights the 2x2 submatrix starting at row 2, column 2. Blue arrows point from the corners of the yellow square to the corners of the teal square.

1	1	1
1	1	



# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{array}{|c|} \hline -1 \\ \hline \end{array} \times \begin{array}{|c|} \hline -1 \\ \hline \end{array} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

1	1	1
1	1	1
1	1	1

$$\frac{1+1+1+1+1+1+1+1+1}{9} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$-1 \times 1 = -1$

The grid is a 9x9 matrix of values -1 and 1. The values are arranged in a checkerboard pattern. The grid is as follows:

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

A 3x3x3 neighborhood is highlighted with a green border, centered at (4,4,4). Within this neighborhood, a 1x1x1 element at (4,4,4) is highlighted with a yellow border. A blue line indicates the path of the sliding window across the grid.

1	1	-1





# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

1	1	-1
1	1	1
-1	1	1

$$\frac{1+1-1+1+1+1-1+1+1}{9} = .55$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

The diagram shows a large rectangle on a grid. A vertical line divides the rectangle into a square on the left and a smaller rectangle on the right. The square is labeled '1' and the smaller rectangle is labeled '.55'. This represents the first step of the Euclidean algorithm, where the larger number (1) is divided by the smaller number (0.55), resulting in a quotient of 1 and a remainder of 0.55.

# Convolution: Trying every possible match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# Convolution: Trying every possible match

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	1
-1	1	-1
1	-1	1

=

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1




-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# Convolution layer

One image becomes a stack of filtered images



# Convolution layer

One image becomes a stack of filtered images

-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1
-1	-1	1	-1	-1	-1	1	-1
-1	-1	-1	1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1
-1	1	-1	-1	-1	-1	1	-1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# Pooling: Shrinking the image stack

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.

# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00			



# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	0.33		

# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	0.33	0.55	


# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33	
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11	
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55	
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33	
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11	
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11	
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77	

1.00	0.33	0.55	0.33

# Pooling



# Pooling

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

max pooling

1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33



0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33


# Pooling layer

A stack of images becomes a stack of smaller images.

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

# Normalization

- Keep the math from breaking by tweaking each of the values just a bit.
- Change everything negative to zero.
- Using activation function.



# Activation Function


- Defines the output of a node.
- Produce a non-linear decision boundary via linear combinations.

# Different Types Of Activation Function

## Activation Functions


### Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




### tanh

$$\tanh(x)$$




### ReLU

$$\max(0, x)$$



### Leaky ReLU

$$\max(0.1x, x)$$




### Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

### ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$





# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)




# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)







0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77


# ReLU layer

A stack of images becomes a stack of images with no negative values.



# Layers get stacked


The output of one becomes the input of the next.






# Deep stacking

Layers can be repeated several (or many) times.




# Fully connected layer

Every value gets a vote




# Fully connected layer

Vote depends on how strongly a value predicts X or O




# Fully connected layer

Vote depends on how strongly a value predicts X or O




# Fully connected layer

Future values vote on X or O




# Fully connected layer

Future values vote on X or O




# Fully connected layer

Future values vote on X or O



# Fully connected layer


Future values vote on X or O






# Fully connected layer

Future values vote on X or O




# Fully connected layer

Future values vote on X or O



# Fully connected layer




A list of feature values becomes a list of votes.



# Fully connected layer


These can also be stacked.

0.9
0.65
0.45
0.87
0.96
0.73
0.23
0.63
0.44
0.89
0.94
0.53




# How neural networks work

# A four pixel camera



# Categorize images



solid



vertical




diagonal



horizontal



# Categorize images



solid



vertical



diagonal




horizontal





# Categorize images



solid



vertical




diagonal



horizontal



# Categorize images



solid



vertical




diagonal



horizontal



# Simple rules can't do it



solid



vertical





diagonal



horizontal



# Simple rules can't do it



solid




vertical




diagonal




horizontal




# Input neurons




# Pixel brightness



# Input vector




# Receptive fields






# A neuron




# Sum all the inputs




$$\begin{array}{r} .50 \\ 0.00 \\ -.75 \\ + \quad .75 \\ \hline .50 \end{array}$$

# Weights




$$\begin{array}{r} .50 \times 1.0 \\ 0.00 \times 1.0 \\ -.75 \times 1.0 \\ + .75 \times 1.0 \\ \hline .50 \end{array}$$

# Weights




$$\begin{array}{rcl} .50 & \times & -.2 \\ 0.00 & \times & 0.0 \\ -.75 & \times & .8 \\ + & .75 & \times & -.5 \\ \hline & & -1.075 \end{array}$$

# Weights




$$\begin{array}{rcl} .50 & \times & -.2 \\ 0.00 & \times & 0.0 \\ -.75 & \times & .8 \\ + & .75 & \times & -.5 \\ \hline & & -1.075 \end{array}$$


# Squash the result



# Sigmoid squashing function




# Sigmoid squashing function







# Sigmoid squashing function




# Sigmoid squashing function




# Sigmoid squashing function




# Sigmoid squashing function




No matter what you start with, the answer stays between  $-1$  and  $1$ .




# Squash the result



# Weighted sum-and-squash neuron




# Make lots of neurons, identical except for weights




To keep our picture clear, weights will either be  
1.0 (white)  
-1.0 (black) or  
0.0 (missing)




# Receptive fields get more complex




# Repeat for additional layers




# Receptive fields get still more complex




# Repeat with a variation




# Rectified linear units (ReLU)




If your number is positive, keep it.  
Otherwise you get a zero.





# Add an output layer























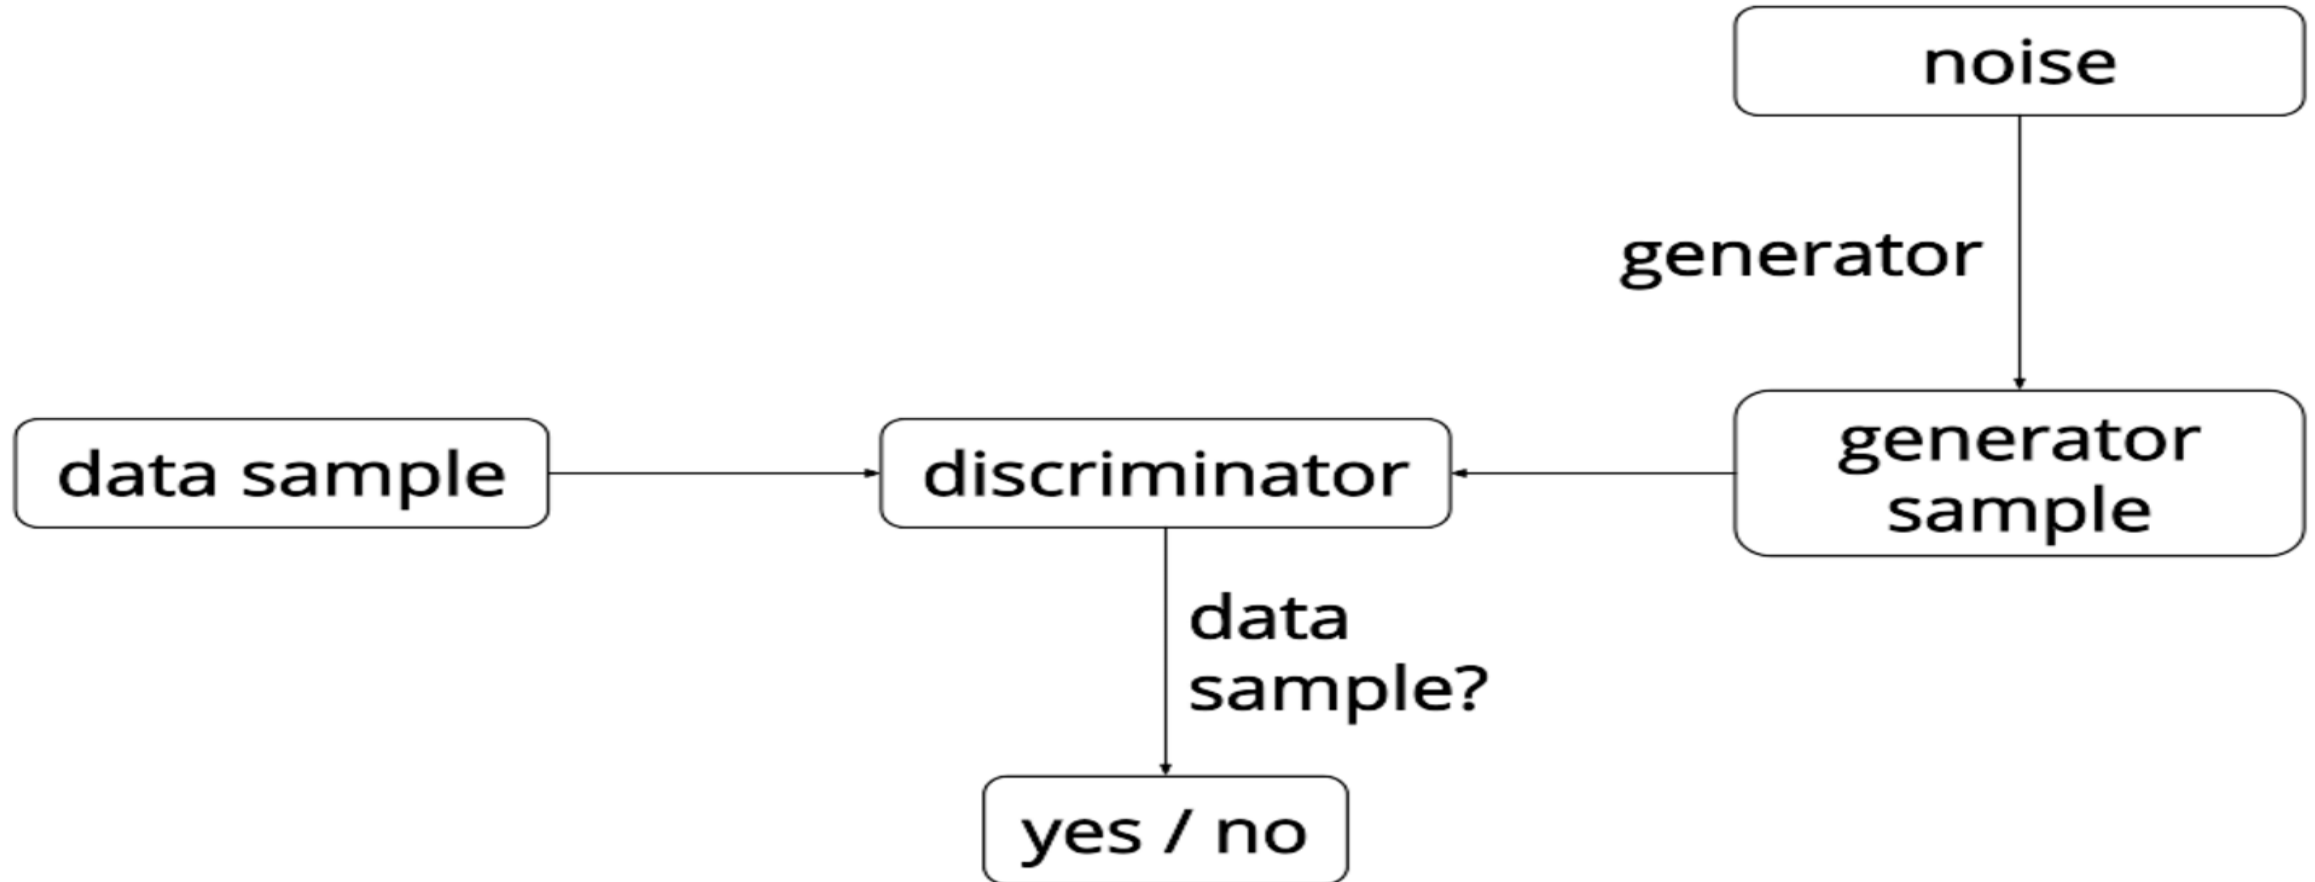
# Generative Adversarial Network(GAN) Architecture

# What Is GAN

- System of two neural networks competing against each other in a zero-sum game framework.
- Can learn to draw samples from a model that is similar to data that we give them.



# Overview Of GAN



# Discriminative Models

- A **discriminative** model learns a function that maps the input data ( $x$ ) to some desired output class label ( $y$ ).
- In probabilistic terms, they directly learn the conditional distribution  $P(y/x)$ .

# Generative Models

- A **generative** model tries to learn the joint probability of the input data and labels simultaneously i.e.  $P(x,y)$ .
- Potential to understand and explain the underlying structure of the input data even when there are no labels

# Training GAN

- Objective of generative network – increase the error rate of the discriminative network.
- Objective of discriminative network – decrease binary classification loss.
- Discriminator training – backprop from a binary classification loss.
- Generator training – backprop the negation of the binary classification loss of the discriminator.

# Loss Function

$$\mathcal{L}(\hat{x}) = \min_{x \in data} (x - \hat{x})^2$$

Generator

$$D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$$

Discriminator

## Alternate view of GANs

$$V(D, G) = \min_G \max_D \mathbb{E}_{x \sim p(x)} [\log D(x)] + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

$$D^* = \operatorname{argmax}_D V(D, G)$$

$$G^* = \operatorname{argmin}_G V(D, G)$$

- In this formulation, Discriminator's strategy was  $D(x) \rightarrow 1$ ,  $D(G(z)) \rightarrow 0$

- Alternatively, we can flip the binary classification labels i.e. **Fake = 1**, **Real = 0**

$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log(1 - D(x))] + \mathbb{E}_{z \sim q(z)} [\log(D(G(z)))]$$

- In this new formulation, Discriminator's strategy will be  $D(x) \rightarrow 0$ ,  $D(G(z)) \rightarrow 1$

## Alternate view of GANs (Contd.)

- If all we want to encode is  $D(x) \rightarrow 0, D(G(z)) \rightarrow 1$

$$D^* = \operatorname{argmax}_D \mathbb{E}_{x \sim p(x)} [\log(1 - D(x))] + \mathbb{E}_{z \sim q(z)} [\log(D(G(z)))]$$

We can use this

$$D^* = \operatorname{argmin}_D \mathbb{E}_{x \sim p(x)} \log(D(x)) + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

- Now, we can replace cross-entropy with any loss function (**Hinge Loss**)

$$D^* = \operatorname{argmin}_D \mathbb{E}_{x \sim p(x)} D(x) + \mathbb{E}_{z \sim q(z)} \max(0, m - D(G(z)))$$

- And thus, instead of outputting probabilities, Discriminator just has to output:-
  - High values for fake samples
  - Low values for real samples

# Generator Network Architecture




Figure : Architecture of Generator and Discriminator Network with corresponding kernel size (k), number of feature maps (n) and stride (s) indicated for each convolutional layer.



# Discriminator Network Architecture

## Discriminator Network




Figure : Architecture of Generator and Discriminator Network with corresponding kernel size (k), number of feature maps (n) and stride (s) indicated for each convolutional layer.

# Perceptual Loss Function

$$l^{SR} = \underbrace{l_X^{SR}}_{\text{content loss}} + \underbrace{10^{-3} l_{Gen}^{SR}}_{\text{adversarial loss}}$$

perceptual loss (for VGG based content losses)

# Content Loss

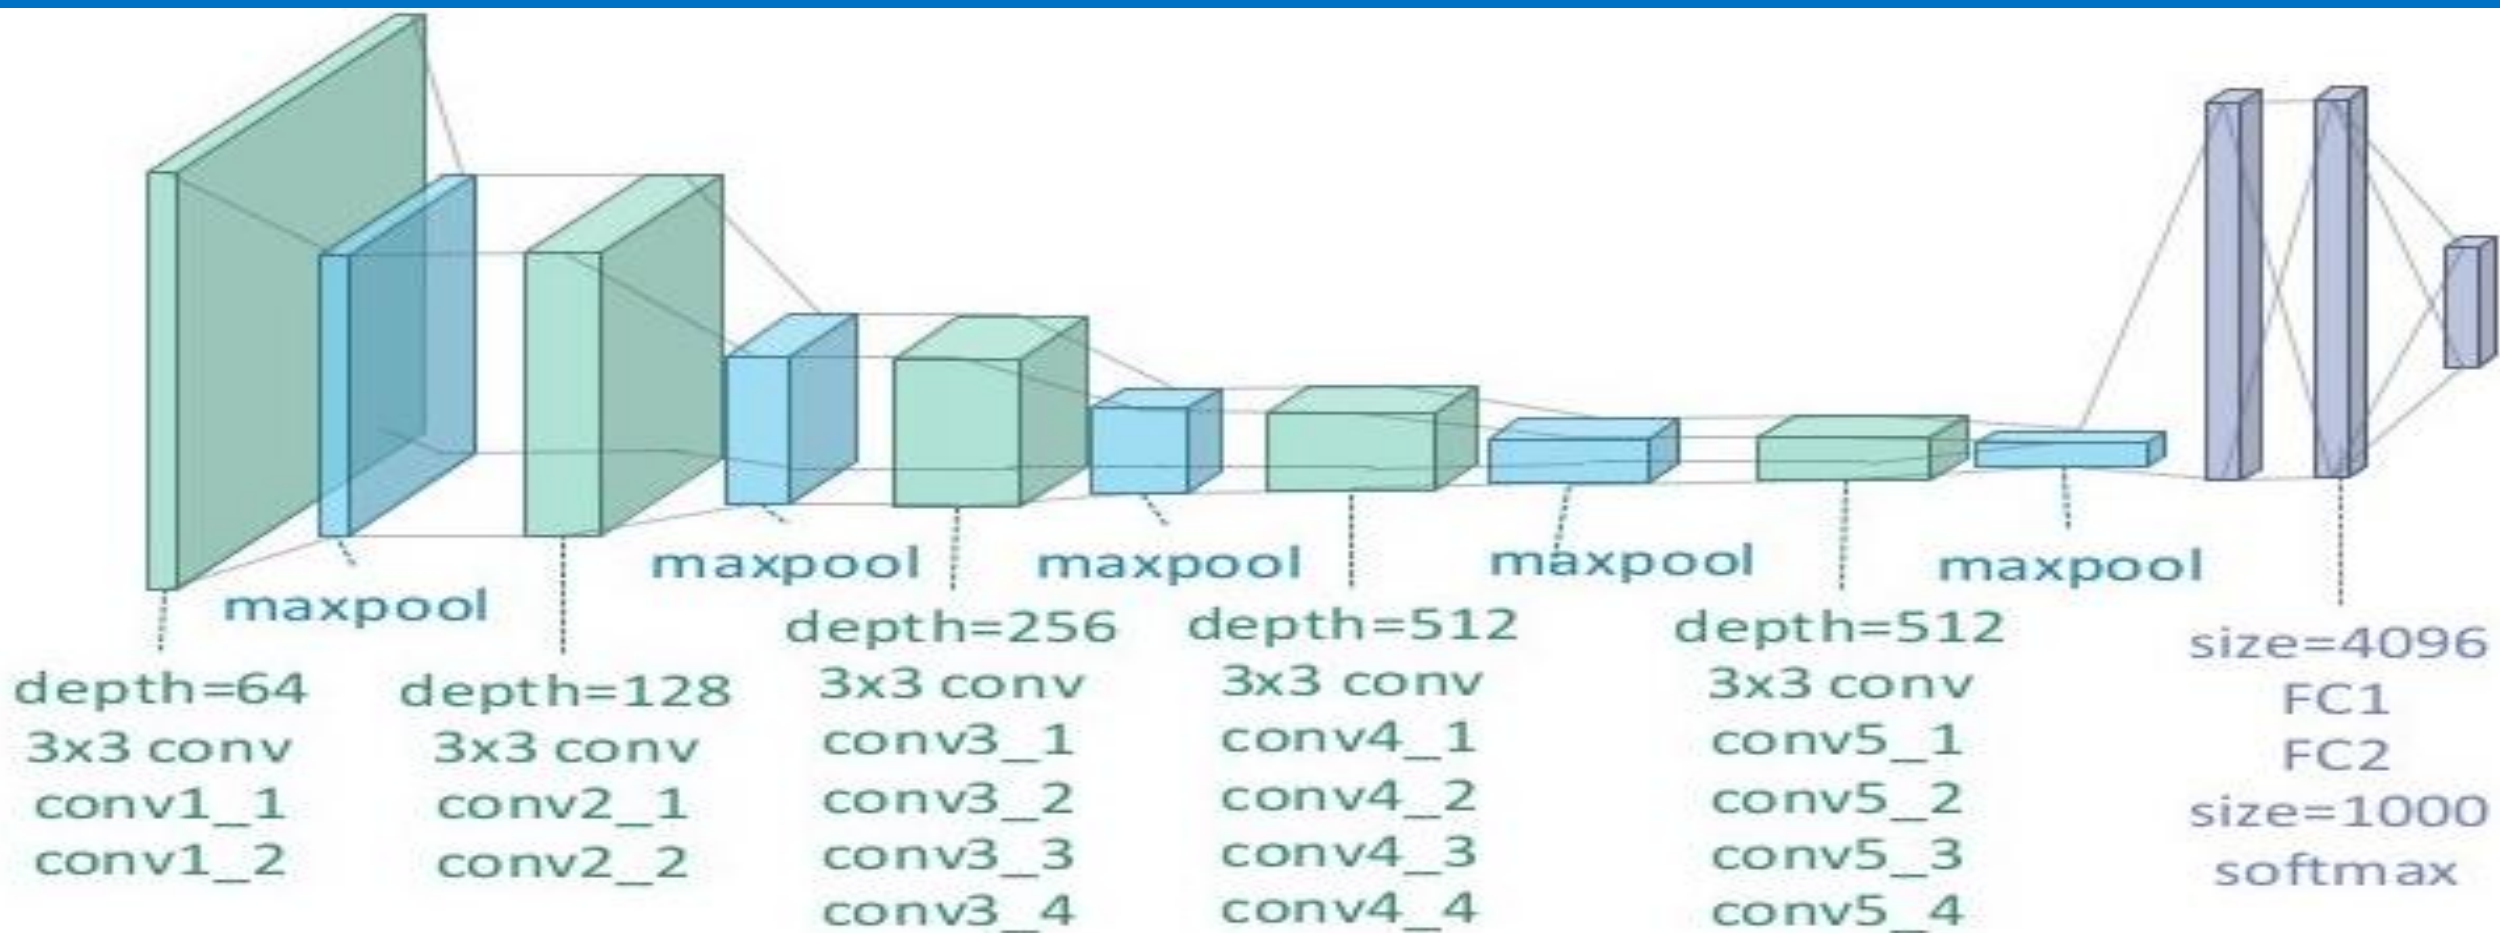
$$l_{MSE}^{SR} = \frac{1}{r^2 W H} \sum_{x=1}^{rW} \sum_{y=1}^{rH} (I_{x,y}^{HR} - G_{\theta_G}(I^{LR})_{x,y})^2$$

$$l_{VGG/i,j}^{SR} = \frac{1}{W_{i,j} H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$$

# VGG19

- Pretrained VGG-19 convolutional neural network
- Trained on more than a million images from the ImageNet database.
- The network is 19 layers deep and can classify images into 1000 object categories

# VGG19 Model



# Adversarial Loss

$$l_{Gen}^{SR} = \sum_{n=1}^N -\log D_{\theta_D}(G_{\theta_G}(I^{LR}))$$

# Peak Signal To Noise Ratio

- The ratio between the maximum possible power of an image and a power of corrupting noise.
- Compare image to an “ideal” clean image with the maximum possible power

# Peak Signal To Noise Ratio

$$PSNR = 10 \log_{10} \left( \frac{(L-1)^2}{MSE} \right) = 20 \log_{10} \left( \frac{L-1}{RMSE} \right)$$



# Mean Squared Error

$$MSE = \frac{1}{NM} \sum_{i=0}^{N-1} \sum_{j=0}^{M-1} (A(i, j) - B(i, j))^2$$

$$RMSE = \sqrt{MSE} = \sqrt{\frac{1}{NM} \sum_{i=0}^{N-1} \sum_{j=0}^{M-1} (A(i, j) - B(i, j))^2}$$

# Peak Signal To Noise Ratio


- PSNR <30.0 (this corresponds to RMSE >11) is commonly considered low. This means the presence of clearly visible noise or smoothing of many edges
- PSNR > 30.0 (this corresponds to RMSE <8.6) is commonly considered as acceptable (some noise is still visible or small details are still smoothed)
- PSNR > 33.0 (this corresponds to RMSE <5.6) is commonly considered as good
- PSNR > 35.0 (this correspond to RMSE <4.5) is commonly considered as excellent (it is usually not possible to find any visual distinction from the “ideal” image)

# Related Work



# Bicubic

- Image resampling.
- Applying linear interpolation in two directions.
- Different from bi-linear interpolation.

# Bicubic



# Bicubic





10	20
30	40

$\xrightarrow{2x}$

P1	P2		

10	10	20	20
10	10	20	20
30	30	40	40
30	30	40	40

# Bicubic



# Neighborhood embedding approach

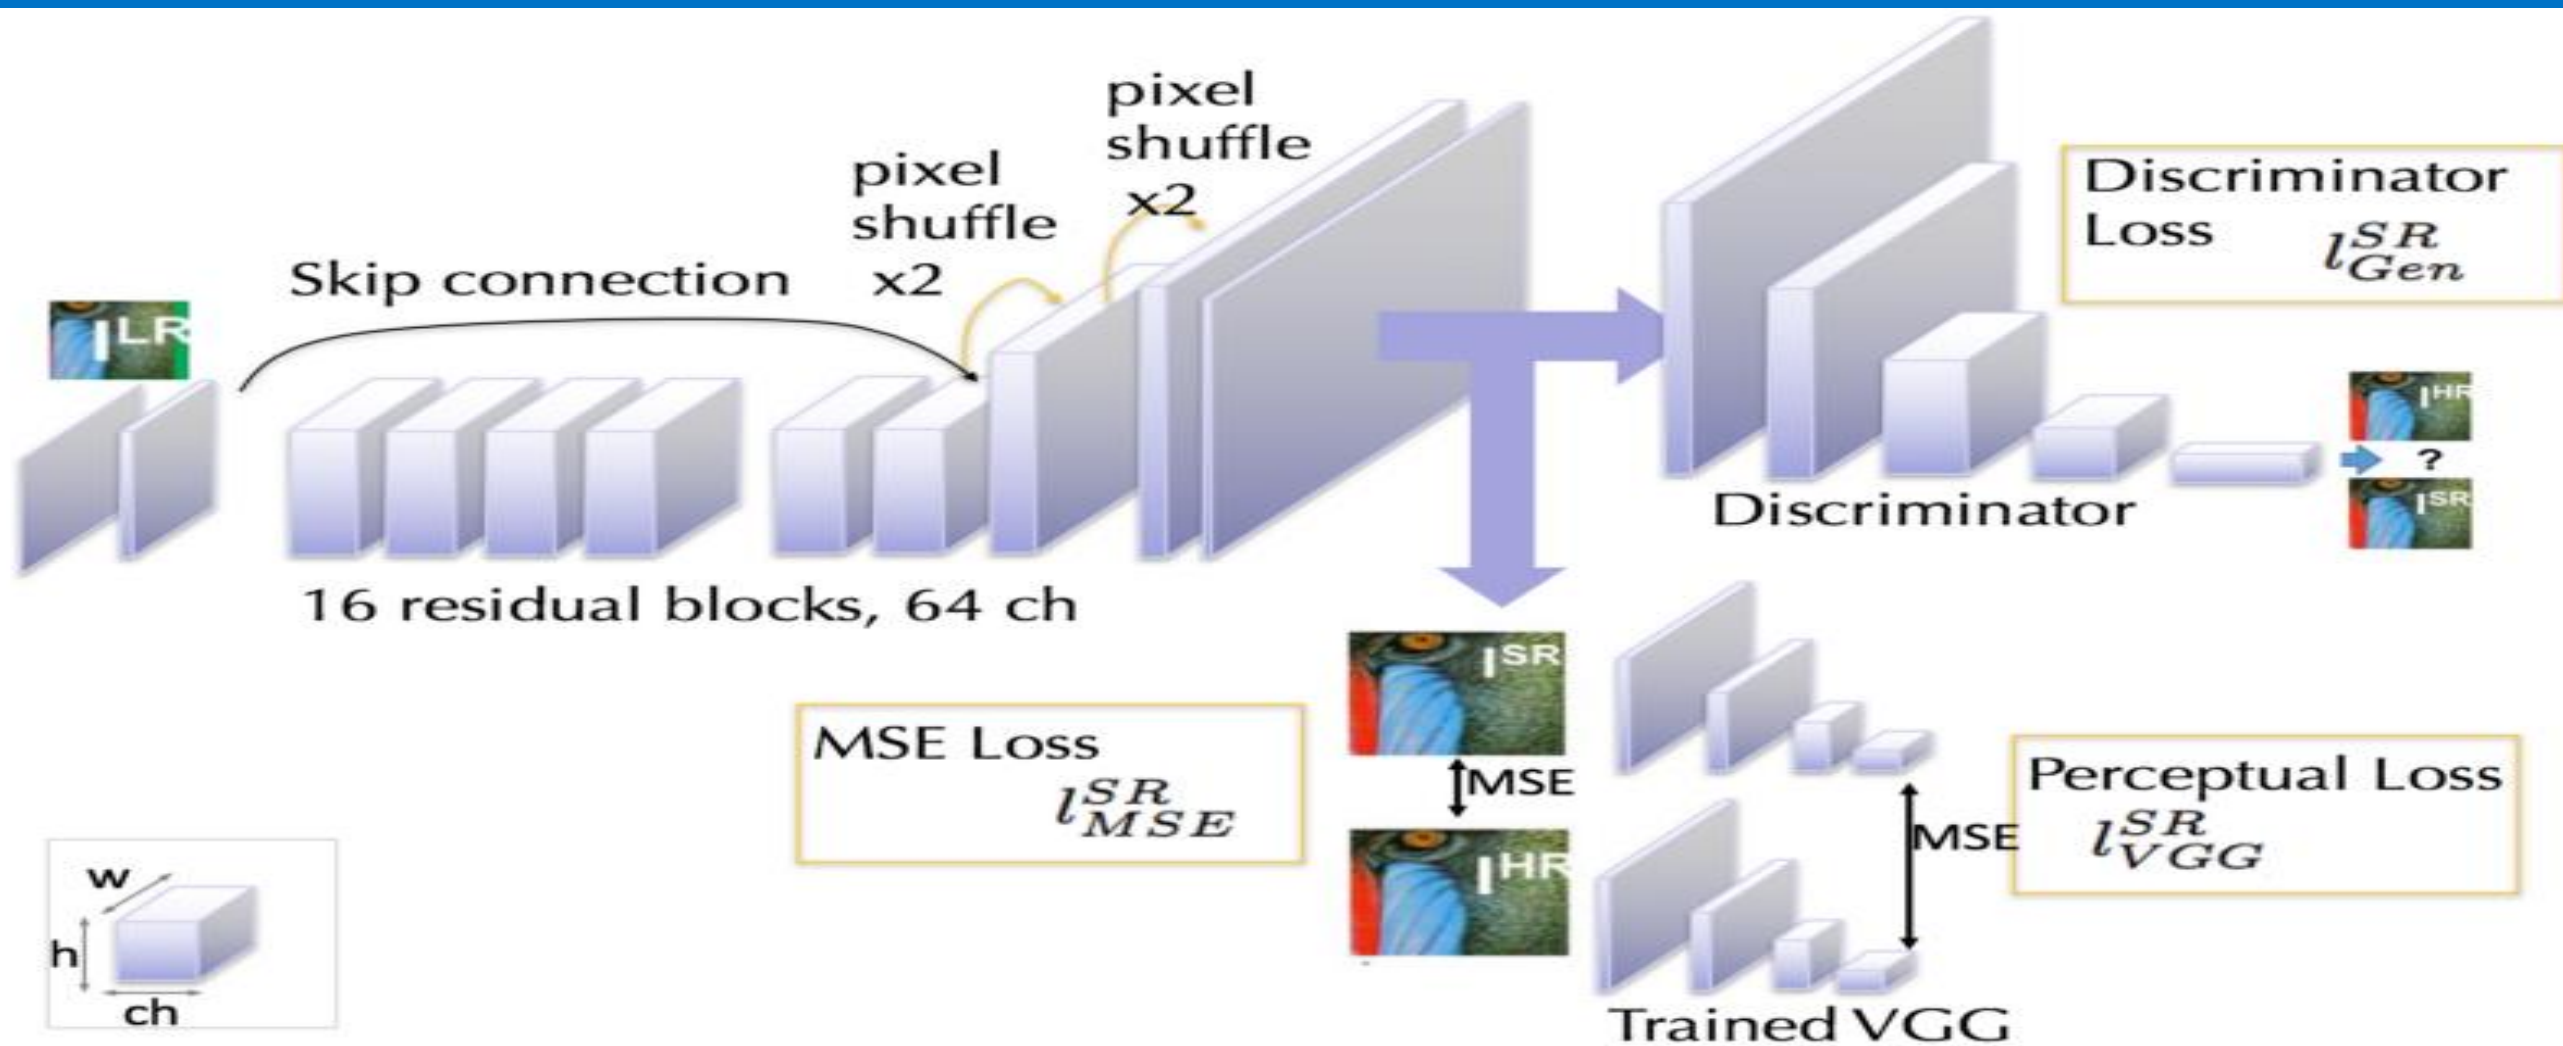
- Uses reconstruction methods.
- Upsample lower resolution image patch
- Uses kernel ridge regression.



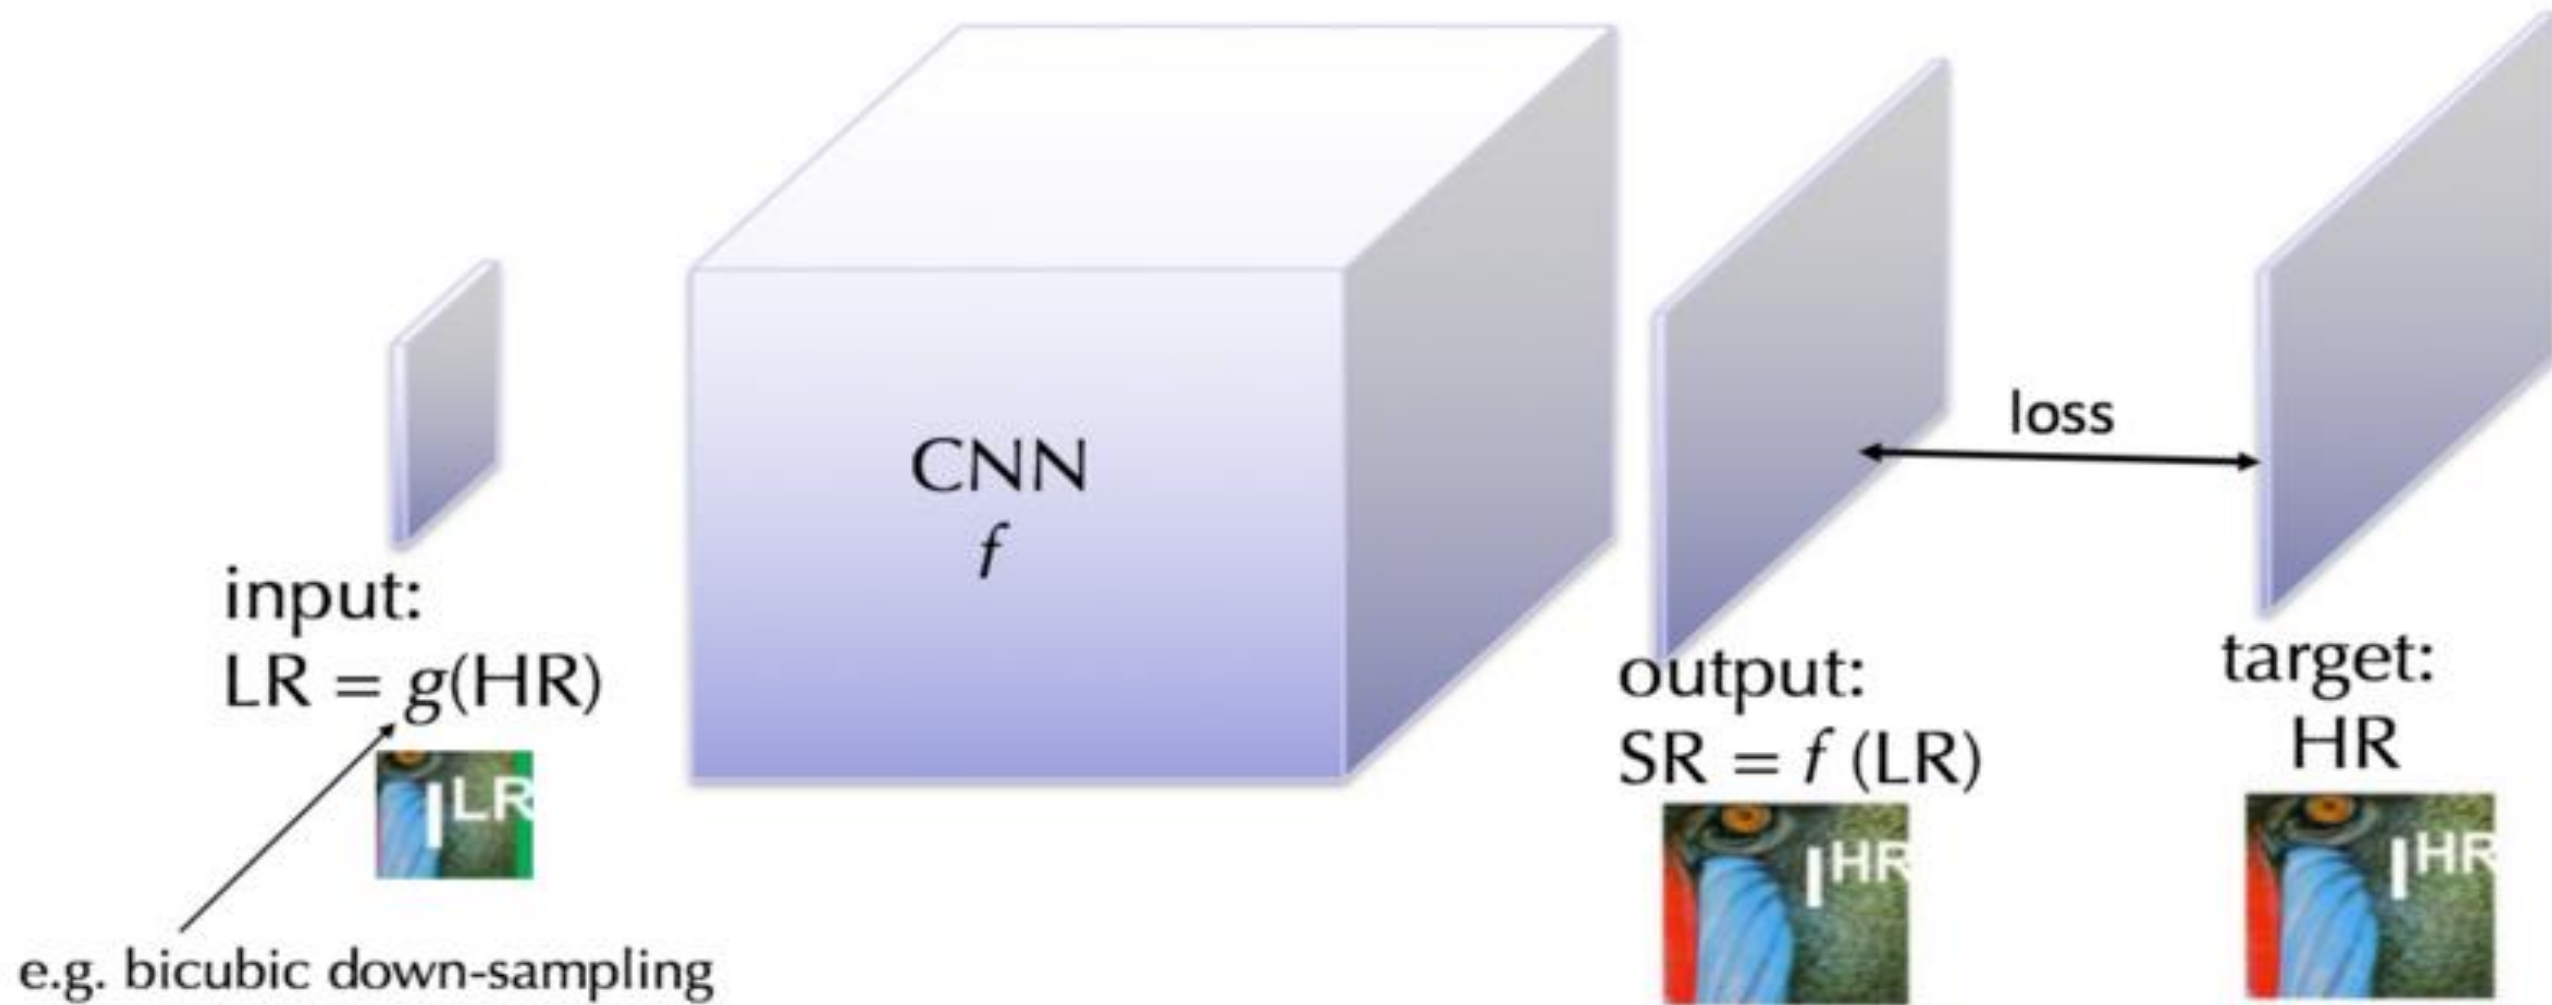
# SRResNet

- Based on ResNet architecture.
- Series of resblocks.
- For upsampling, pixel shuffle operator is used.


# SRResNet



# Concept of SISR



# Results



# Comparison Table

Set5	nearest	bicubic	SRCNN	SelfExSR	DRCN	ESPCN	SRResNet	SRGAN	HR
PSNR	26.26	28.43	30.07	30.33	31.52	30.76	<b>32.05</b>	29.40	$\infty$
SSIM	0.7552	0.8211	0.8627	0.872	0.8938	0.8784	<b>0.9019</b>	0.8472	1
MOS	1.28	1.97	2.57	2.65	3.26	2.89	3.37	<b>3.58</b>	4.32
Set14									
PSNR	24.64	25.99	27.18	27.45	28.02	27.66	<b>28.49</b>	26.02	$\infty$
SSIM	0.7100	0.7486	0.7861	0.7972	0.8074	0.8004	<b>0.8184</b>	0.7397	1
MOS	1.20	1.80	2.26	2.34	2.84	2.52	2.98	<b>3.72</b>	4.32
BSD100									
PSNR	25.02	25.94	26.68	26.83	27.21	27.02	<b>27.58</b>	25.16	$\infty$
SSIM	0.6606	0.6935	0.7291	0.7387	0.7493	0.7442	<b>0.7620</b>	0.6688	1
MOS	1.11	1.47	1.87	1.89	2.12	2.01	2.29	<b>3.56</b>	4.46

